

Analyzing Toronto Apartment Rental Prices

Austin Hart (45726379)

Olivia Jiang (82785783)

Woobin Kim (39017496)

Raina Tian (11946720)

2026-07-27

Project Task Assignments:

- Austin did...
- Olivia did...
- Raina did...
- Woobin did...

Introduction

In high-density metropolitan areas like Toronto, assessing how specific layout configurations influence rental pricing provides valuable insight into market structure and housing affordability. This study examines the relationship between apartment layout characteristics and monthly rent across the city. Specifically, we use multiple linear regression to evaluate how the number of bedrooms, bathrooms, and their interaction effects predict continuous rental prices. To test predictive validity further, we apply binary logistic regression using the same structural features to determine whether a listing classifies as a high-value property. By comparing main-effects and interaction models across both regression frameworks, we assess the overall explanatory power of layout features in Toronto's housing market.

Data Description

We obtained the data from [this](#) Kaggle dataset. The amount of rows in the original dataset before any data cleaning is 1,124. The data contains information on apartment rental units in Toronto collected in 2018, including the apartment's geographic information (**address**, **latitude**, and **longitude**), information on its rooms (the number of **bedrooms**, **bathrooms**, and **dens**), and the unit's monthly rental **price**.

The total number of rows in the original dataset is 1,124. There are no missing values in any columns.

During data cleaning, we removed 308 duplicates which likely occurred because the data came from multiple sources. We also removed 5 rows with unusual values for **price** (\$65/month, \$99/month, \$99/month, \$36,900/month, \$535,000/month). These outliers may skew our results so we decided to remove them. The sample size is now **n = 811**.

Exploratory Data Analysis

Before modeling, we want to find out more about the data.


We first examine each variable individually. Below are the sample's summary statistics of each variable.

Variable	Mean	SD	Min	Max
bathroom	1.24	0.44	1.00	3.00
bedroom	1.37	0.57	1.00	3.00
den	0.15	0.36	0.00	1.00
lat	43.71	0.69	42.99	56.13
long	-79.50	1.84	-114.08	-73.58
price	2185.71	787.36	150.00	9750.00

TODO: commentary on meaningfulness... There's nothing unusual... As well, **price**'s standard deviation is \$1,021/month, its minimum is \$150/month, and its maximum is \$9,750/month. We decided that...

We also visualize the distribution of each variable individually.

To visualize the distribution of `price`, we plot a histogram.




The distribution is slightly right-skewed and unimodal around the mean of `mean`. As well, its standard deviation is \$1,021/month, the minimum is \$150/month, and the maximum is \$9,750/month.

TODO: more commentary

Next, we view the frequencies of each discrete value of `bedroom`, `bathroom`, and `den`.


Frequencies of Room Counts by Room Type



TODO: comment on them (such as the unbalanced groups)

Next, we examine possible relationships between variables. Below are boxplots of `price` split by `bedroom`, `bathroom`, and `den` counts in order to see how `price`'s distribution varies by each of these 3 apartment characteristics.


Distributions of Rent Prices by Room Counts



TODO: comment on them

TODO: number of observations of each 0/1 in `high_price`

We're also curious about the correlations between the variables so we plot (the upper triangle of) a correlation matrix.



The variables with the strongest correlation with `price` are `bathroom` ($r = 0.47$) and `bedroom` ($r = 0.38$), though these are only moderately strong linear correlations. `den` is weakly negatively correlated with `price` at $r = -0.01$. Moreover, `bedroom` and `bathroom` are moderately correlated at $r = 0.50$, which may indicate that they're collinear, which may affect our modeling decisions later on.

TODO: key take-aways of EDA

Models

Now that we have processed our data, we can consider models that can predict rental price based on our other variables. In the previous section, we narrowed down the variables we can use as predictors to number of bedrooms and number of bathrooms. To make use of these predictors, we have two options: a multiple linear regression and a multiple logistic regression.

Multiple Linear Regression

A multiple linear regression model generates predictions as a straight line with a slope and intercept.

When making a multiple linear regression model, one of the first things to consider is whether or not interaction exists between covariates. If there is interaction, then the slope of the prediction line will be different depending on the value of the covariates; if not the slope is constant. This can be checked with hypothesis testing, something which can be checked by creating a model with interaction in R and checking the model's statistics.

Results with interaction:

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	1609.69	238.10	6.76	0.00	1142.32	2077.07
bedroom	-131.33	141.90	-0.93	0.35	-409.86	147.20
bathroom	215.80	203.38	1.06	0.29	-183.42	615.02
bedroom:bathroom	263.75	102.30	2.58	0.01	62.94	464.56


Results without interaction:

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	1026.44	74.50	13.78	0	880.20	1172.68
bedroom	206.60	54.54	3.79	0	99.53	313.66
bathroom	707.66	70.73	10.01	0	568.83	846.49


As shown in the tables above, the table produced by the model with interaction is reporting p-values in excess of our predecided level of significance (95%, necessitating a value below 0.05). Thus, we cannot confidently identify any interaction between the predictors and must proceed with the model without interaction.

The estimate column in the table for the non-interaction model shows that our prediction line will have an intercept of 1026.44. It shows that an additional bathroom in a property is associated with a \$206.60 increase in monthly rent when number of bathrooms is held constant and that an additional bathroom in a property is associated with a \$707.66 increase in rent when number of bedrooms is held constant.

To ensure that our linear model is appropriate for the data, we need to run some model diagnostics in the form of a residual and Q-Q plot.



For a linear model to be appropriate, the error around the trend/prediction line should be consistent and normally distributed. Residual plots offer a convenient visual representation of this error. In a residual plot, error should be distributed normally and evenly around a horizontal line at $y = 0$. As you can see in the plot above, because our predictors are categorical, the data points are not continuously distributed, making the plot somewhat messy to look at. It shows that while the error approximately follows a normal distribution, but with larger variance around certain areas in the x axis, meaning the variance is not perfectly consistent.



Another way of checking the distribution of the error is a Q-Q plot. If the error is normally distributed with equal variance, then a Q-Q plot will show all data points distributed along a straight, sloped line wherein $X = Y$. The Q-Q plot of our linear model shows the data following this line for the most part, but diverging at the beginning and the end, implying that our model has more extreme data than can be expected under a true normal distribution (Clay, 2015). Applying transformations to the model, such as log transformation and predicting the square root of the response variable, were unsuccessful in creating a more favourable distribution.

While our diagnostics reveal that our data may not be perfectly suited for a linear analysis, this is not uncommon for observed data. It is our belief that our model is still useful for making predictions.

Multiple Logistic Regression

A multiple logistic regression model is used to make predictions for a categorical response variable. To make use of such a model, we converted our response variable (price) into a categorical one by splitting it into two categories for rent above and below the mean value. Thus, our model will be predicting whether a property has above or below average rent, rather than a specific value.

Once again, we need to begin by checking for interaction between predictors.

Results with interaction:

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	-0.29	0.16	-1.77	0.08	-0.61	0.03
bedroom	0.19	0.10	1.96	0.05	0.00	0.38
bathroom	0.57	0.14	4.07	0.00	0.29	0.84
bedroom:bathroom	-0.11	0.07	-1.58	0.11	-0.25	0.03

Results without interaction:

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	-0.04	0.05	-0.85	0.40	-0.14	0.06
bedroom	0.05	0.04	1.30	0.19	-0.02	0.12
bathroom	0.36	0.05	7.44	0.00	0.27	0.46

The tables above show that neither model passes the hypothesis tests necessary to prove statistical significance. Of the two models, the results with interaction come much closer. To proceed with our analysis, our only option is to lower our stated level of significance to 89% and continue with the model with interaction. So, keep in mind that any further observations on our logistic regression model will be made with this altered level of significance.

The table of our model with interaction shows that an additional bedroom increases the log odds of a property having high rent by 0.19, an additional bathroom increases the log odds of a property having high rent by 0.57, and .

Again we must run some model diagnostics to ensure our model is appropriate for the data. This time we will use a binned residuals plot.

(insert binned residuals plot here).

In a binned residuals plot, the majority of data points (approximately 95%) must fall within the generated error bounds (Ludecke et al., n.d.). As you can see, our model easily meets this criteria as all of the data points fall within these bounds.

Conclusion

This section should provide overall conclusions, any limitations in the datasets and analyses, and validity of the conclusions.

Appendix

```
# Loading packages  
library(tidyverse)  
library(knitr)  
library(scales)  
library(patchwork)  
library(corrplot)  
library(broom)
```