

Analyzing Toronto Apartment Rental Prices

Austin Hart (45726379)

Olivia Jiang (82785783)

Woobin Kim (39017496)

Raina Tian (11946720)

August 2, 2026

Project Task Assignments

Raina wrote most of the introduction and data description.

Olivia (group leader) did most of the exploratory data analysis section.

Austin did most of the modeling section.

Woobin wrote most of the conclusion.

All members contributed to editing and proofreading the project as a whole.

Introduction

In high-density metropolitan areas like Toronto, assessing how specific layout configurations influence rental pricing provides valuable insight into market structure and housing affordability. This study examines the relationship between apartment layout characteristics and monthly rent across the city. Specifically, we use multiple linear regression to evaluate how the number of bedrooms, bathrooms, and their interaction effects predict continuous rental prices. To test predictive validity further, we apply binary logistic regression using the same structural features to determine whether a listing classifies as a high-value property. By comparing main-effects and interaction models across both regression frameworks, we assess the overall explanatory power of layout features in Toronto's 2018 housing market.

Data Description

The dataset for this study was retrieved from a public Kaggle repository titled Toronto Apartment Rental Price, compiled by independent creator Raja CSP. According to the dataset documentation, the listings were aggregated and scraped from local rental websites active across Toronto in 2018.

The study utilizes Version 3 of the dataset, stored as a single 94.97 KB comma-separated values file. The raw file contains $n = 1,124$ observational rows across 7 attribute columns, where each row corresponds to a unique apartment listing. There are no missing values in any columns in the raw data.

Exploratory Data Analysis

During data cleaning, we removed 308 duplicates which likely occurred because the data came from multiple sources. We also removed 9 rows where **price** was either below \$100/month or above \$6,000/month (namely, \$65, \$99, \$99, \$6,000, \$6,500, \$8,000, \$9,750, \$36,900, and \$535,000) because these outliers may skew our results. The sample size is now $n = 807$.

Before modeling, we want to find out more about the data.

We first examine sample's summary statistics of each individual numerical variable.


Variable	Mean	Standard Deviation	Minimum	Maximum
bathroom	1.23	0.43	1.00	3.00
bedroom	1.36	0.56	1.00	3.00
den	0.15	0.36	0.00	1.00
lat	43.71	0.69	42.99	56.13
long	-79.50	1.85	-114.08	-73.58
price	2159.05	684.27	150.00	4900.00

The amount of **bathrooms** ranges from 1 to 3, the amount of **bedrooms** ranges from 1 to 3, and the amount of **dens** ranges from 0 to 1. The average amount of bedrooms is slightly higher than the average amount of bathrooms, and most apartment's don't have any dens, shown by its mean of 0.15. This makes sense - an apartment should have a bedroom and bathroom whereas a den is optional. Also, **latitude** and **longitude** fall in the valid range of $[-90, 90]$ and $[-180, 180]$ respectively (though we won't use these variables in modeling). So there's nothing unusual about these four variables.

As for **price**, its average value is \$2,159/month with a standard deviation of \$684/month. Its minimum is \$150/month and its maximum is \$4,900/month. While these minimum and maximum values are -2.59 and 4.01 standard deviations above and below the mean (respectively), we feel that they are reasonable enough to keep.

We now plot a histogram to visualize the distribution of `price` and a bar chart showing the class balance of the new binary variable `high_price`.

Distributions of Toronto Monthly Apartment Rent and High Price Indicator (2018)




On the left, we see that `price` is approximately symmetric and unimodal around its mean. On the right, we see that the two `high_price` groups are approximately balanced; this means that logistic regression won't run into the issue of unbalanced response groups, which could bias modeling results.

Next, we examine possible relationships *between* variables.


Below are boxplots of `price` split by `bedroom`, `bathroom`, and `den` counts to see how `price`'s distribution varies by each of these apartment characteristics.

Distributions of Toronto Rent Prices by Room Counts (2018)



Looking independently at `bedroom`, `price`'s distribution slightly shifts upwards by a few hundred dollars for each additional bedroom. For `bathroom`, the distribution also moves upwards with the biggest drop from 2 to 2.5 bathrooms and another upwards shift in the distribution from 2.5 to 3 bathrooms. For `den`, the distribution is roughly equal for both groups with a difference-between-means of only \$87/month. As well, within each group of boxplots, the distributions overlap with each other, so these predictors may not produce strong modeling results. Nonetheless, in terms of potential modeling power, `bathroom` may be better than `bedroom` and it may be better not to use `den`.

We're also curious about the correlations between the variables so we plot the upper triangle of a (Pearson) correlation matrix. Pearson correlation, denoted by the letter r , falls in $[0, 1]$. If two variables are positively correlated, it means that as one gets larger, the other also tends to get larger; if two variables are negatively correlated, it means that as one gets larger, the other tends to get smaller.



Looking at the rightmost column, the variables with the strongest linear correlation with **price** are **bathroom** ($r = 0.47$) then **bedroom** ($r = 0.39$), though these are only moderately strong. In addition to the boxplots' evidence, the matrix further confirms that **den** is weakly correlated with **price** at $r = 0.05$. And even though **bedroom** and **bathroom** are moderately correlated at $r = 0.63$, putting models at risk of multicollinearity, both variables measure different characteristics of an apartment, so we may keep both during modeling.

In addition, here are frequency tables between **high_price** and **bedroom** and **bathroom**, excluding **den** because we've already seen that it would be a weak predictor.

# of Bedrooms	Rent Below Average	Rent Above Average	Total
1	62% (340)	38% (208)	100% (548)
2	32% (72)	68% (153)	100% (225)
3	38% (13)	62% (21)	100% (34)
Total	53% (425)	47% (382)	100% (807)

# of Bathrooms	Rent Below Average	Rent Above Average	Total
1	62% (381)	38% (232)	100% (613)
1.5	52% (12)	48% (11)	100% (23)
2	18% (30)	82% (135)	100% (165)
2.5	100% (1)	0% (0)	100% (1)
3	20% (1)	80% (4)	100% (5)
Total	53% (425)	47% (382)	100% (807)

Most of the data is on 1-bedroom or 1-bathroom units with the least data on units with 1.5, 2.5, or 3 bathrooms. Expectedly, most (62%) of the below-average units have 1 bedroom and 1 bathroom.

The main findings of this exploratory analysis are:

1. There is no class imbalance in **high_price** (for the purpose of logistic regression) because 53% vs. 47% of the observations have a **price** above vs. below average. Class imbalance, such as 90% vs. 5%, could bias a model's intercept and produce a poor model that maximizes accuracy by simply predicting the majority class.
2. In terms of predictive ability, **den** likely has the weakest linear relationship with **price**, as shown by the very similar boxplots and Pearson correlation of 0.05, and we exclude **latitude** and **longitude** because geographic effects on a target variable are probably not as linear or straightforward as these coordinates. This leaves **bedroom** and **bathroom** as the covariates remaining.

Models

Now that we have explored the data, we can consider models that can model rental price based on other variables. In the previous section, we narrowed down the variables we can use as predictors to the number of bedrooms and the number of bathrooms. To make use of these predictors, we chose two model types: a multiple linear regression on `price` because it is continuous and a multiple logistic regression on `high_price` because it is binary.

Multiple Linear Regression

A multiple linear regression model models a continuous response variable in terms of the predictors as a straight line with an intercept and at least one slope.

When making this model type, we often consider whether interaction between covariates is statistically significant. If it is, then the slope of the prediction line differs depending on the value of the covariates, otherwise the slope is constant. Statistical significance can be checked with hypothesis testing by creating a model with interaction and checking its statistics.

Results with interaction:

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	1496.02	210.86	7.09	0.00	1082.12	1909.93
bedroom	0.95	125.45	0.01	0.99	-245.30	247.21
bathroom	318.44	180.91	1.76	0.08	-36.68	673.56
bedroom:bathroom	146.84	90.94	1.61	0.11	-31.67	325.35

Results without interaction:

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	1172.76	66.25	17.70	0.0000000	1042.72	1302.80
bedroom	187.90	48.35	3.89	0.0001101	93.00	282.80
bathroom	592.21	63.16	9.38	0.0000000	468.23	716.19

In the first table, the interaction term (`bedroom:bathroom`) has a p-value of 0.11 in excess of our pre-decided level of significance of 0.05. Thus, we cannot confidently identify an interaction between the predictors and choose to proceed with the model without interaction, which has small p-values for all three coefficients.

The `estimate` column in the table for the non-interaction model shows a prediction line of the form:


$$\text{price} = \$1172.76 + \$187.90 \cdot \text{bedroom} + \$592.21 \cdot \text{bathroom}$$

The intercept value has no obvious interpretation because a place with no bedrooms or bathrooms isn't very realistic; however, the intercept is like a necessary starting point for the prediction so that the slopes stay unbiased.

Both slope values are statistically significant, consistent with the EDA results suggesting that the . The slope values show that each additional bedroom in a property is associated with a \$187.90 increase in monthly rent when the number of bathrooms is fixed at any value, and that an additional bathroom is associated with a \$592.21 increase in rent when the number of bedrooms is fixed at any value. For example, the model predicts that a one-bed, one-bath condo costs $\$1172.76 + \$187.90 + \$592.21 = \$1952.87/\text{month}$.


To check if the fitted model is appropriate, we create visual diagnostics in the form of a residual plot and a Q-Q plot. These essentially check if the assumptions for linear regression to be valid are fulfilled.

A residual plot plots the residuals/errors (the difference between the actual and fitted values) against the fitted values to check for the errors' linearity (i.e., a mean of 0) and constant variance. Errors should be distributed normally and evenly around the $e = 0$ horizontal line.



The error appear to have a mean of 0 because the points appear evenly around both sides of $e = 0$, but the variance is larger at certain areas (like when the fitted values are around 2000 and 3000), so the variance is not perfectly consistent.

A Q-Q plot is another way to check the errors' distribution. If the error is normally distributed, then a Q-Q plot will show all data points distributed along a straight, sloped line wherein $x = y$.



The Q-Q plot of the model shows the data following this line for the most part, but the points diverge at the beginning and the end, implying that the model has more extreme data than can be expected under a true normal distribution (Ford, 2015). Applying transformations to the response variable, such as a log transformation and a square root transformation, were unsuccessful in creating a more favourable distribution.

While our diagnostics reveal that our data may not be perfectly suited for a linear analysis, this is not uncommon for observed data. It is our belief that our model can still be useful for predicting.

Multiple Logistic Regression

A multiple logistic regression model is used to make predictions for a binary response variable so we made **price** binary by splitting it into two categories: above or below the mean. So our model will predict whether a property has above or below average rent, rather than a specific value.

Once again, we begin by checking for interaction between predictors.

Results with interaction:

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	-0.30	0.17	-1.82	0.07	-0.62	0.02
bedroom	0.20	0.10	2.05	0.04	0.01	0.39
bathroom	0.58	0.14	4.13	0.00	0.31	0.86
bedroom:bathroom	-0.12	0.07	-1.70	0.09	-0.26	0.02

Results without interaction:

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	-0.04	0.05	-0.68	0.50	-0.14	0.07
bedroom	0.05	0.04	1.26	0.21	-0.03	0.12
bathroom	0.36	0.05	7.27	0.00	0.26	0.46

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	0.97	0.05	-0.68	0.50	0.87	1.07
bedroom	1.05	0.04	1.26	0.21	0.97	1.13
bathroom	1.43	0.05	7.27	0.00	1.30	1.58

The first table for the non-interaction model models log-odds whereas the second table models odds.

Like our linear model, we cannot conclude that there is significant interaction between the predictors at a 0.05 significance level, so we proceed with the model without interaction.

The **estimate** column in the first table for the non-interaction model shows a prediction line of the form:

$$\log\left(\frac{\hat{p}}{1-\hat{p}}\right) = -0.04 + 0.05 \cdot \text{bedroom} + 0.36 \cdot \text{bathroom}$$

where $\hat{p}$ is the expected probability of a property having high rent and $\frac{\hat{p}}{1-\hat{p}}$ is its odds.


An additional bedroom increases the expected *log-odds* of a property having high rent by 0.05 and increases the expected *odds* of a property having high rent by a factor of $e^{0.05} = 1.05$ or 5%. However, the p-value of 0.21 for this estimate implies that this conclusion is not statistically significant at a 0.05 significance level. Similarly, an additional bathroom increases the odds of a property having high rent by $e^{0.36} = 1.43$ or 43%, which is statistically significant.

Again, we do some model diagnostics.

One of the first things to check for is overdispersion, which occurs when our response distribution deviates from the assumption that the distribution of the fitted values is consistent with a Bernoulli distribution. The model has a dispersion parameter of 0.22, implying insignificant levels of overdispersion.


```
##
## Call:
## glm(formula = high_price ~ bedroom + bathroom, family = "quasibinomial",
##      data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.5505  -0.9601  -0.9601   1.3155   1.4115
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -2.4096     0.2615  -9.216  < 2e-16 ***
## bedroom        0.2164     0.1726   1.254    0.21
## bathroom      1.6578     0.2463   6.730 3.23e-11 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for quasibinomial family taken to be 1.024747)
##
## Null deviance: 1116.4  on 806  degrees of freedom
## Residual deviance: 1013.6  on 804  degrees of freedom
## AIC: NA
##
## Number of Fisher Scoring iterations: 4
```

As an additional assessment, we use a binned residuals plot, where most data points (~95%) should fall within the generated error bounds for the model to be considered appropriate (Lüdecke et al., n.d.).



Each dot is the average residual for one bin of predicted probabilities and the vertical lines show the uncertainty around that average residual. The gray lines show the acceptable range for residuals. Since 2 of the 6 bins fall outside the range, specifically at the estimated probabilities exceeding 80%, the plot shows that the model fits decently well and that higher predicted probabilities should be more carefully interpreted.

Conclusion

The multiple linear regression analysis indicates that both the number of bedrooms and bathrooms have a positive relationship with monthly rent. Specifically, holding other variables constant, an additional bedroom increases predicted monthly rent by \$187.90 ($p < 0.001$), while an additional bathroom increases rent by \$592.21 ($p < 0.001$). The multiple logistic regression analysis indicates that bathrooms have a statistically significant positive effect on the odds of a property having an above-average rent ($p < 0.001$), increasing the odds by 43% per bathroom, while the effect of bedrooms is not statistically significant at the standard 0.05 level ($p = 0.21$).

However, as seen from the residual plot and the Q-Q plot, the key assumptions of homoscedasticity and the normality of residuals for the linear model are not completely valid. These violations affect prediction confidence intervals and the p-values. Namely, the current MLR model results cannot be considered fully reliable. Consequently, findings regarding property characteristics should be interpreted with caution given potential omitted variable bias.

To improve model reliability in future research, two key adjustments are recommended:

- Re-introducing location through neighbourhood clusters may help reduce unexplained variance and standard errors, thereby addressing potential omitted variable bias.
- Collecting continuous structural variables like square footage alongside discrete room counts may help normalize the residual distribution.

Standard 95% confidence thresholds should be maintained to ensure hypothesis testing remains statistically unbiased.

References

- Ford, C. (2016, August 26). *Understanding QQ plots*. University of Virginia Library. <https://library.virginia.edu/data/articles/understanding-q-q-plots>
- Lüdecke, D., Makowski, D., Ben-Shachar, M. S., Patil, I., Waggoner, P., Weirnik, B. M., & Thériault, R. (n.d.). *Binned residuals for binomial logistic regression*. Easystats. Retrieved July 31, 2026, from https://easystats.github.io/performance/reference/binned_residuals.html.

Appendix

```
# Loading packages
library(tidyverse)
library(knitr)
library(scales)
library(patchwork)
library(corrplot)
library(broom)
library(janitor)
library(performance)
library(see)

# Reading the data
data <- read_csv("Toronto_apartment_rentals_2018.csv")

# Seeing the total number of rows
nrow(data)

# Seeing how many duplicate rows there are
sum(duplicated(data)) # 308

# Viewing duplicate rows
data[duplicated(data) | duplicated(data, fromLast = TRUE), ]

# Cleaning the data
data <- data %>%
  # Removing duplicate rows
  distinct() %>%
  # Removing unnecessary columns
  select(-c(Address)) %>%
  # Making the column names lowercase
  rename_with(tolower) %>%
  # Converting price to a number
  mutate(price = parse_number(price)) %>%
  # Removing outliers
  filter(price > 100 & price < 6000)

# Seeing if there are any NAs
colSums(is.na(data)) # 0

# Adding a binary response variable (whether price is more than its mean)
data <- data %>%
  mutate(high_price = if_else(price > mean(price), 1, 0))

# Viewing summary stats
data %>%
  select(-high_price) %>%
  # Reshaping data into a long format
  pivot_longer(cols = everything(), names_to = "Variable", values_to = "Value") %>%
  # Grouping by the variable
  group_by(Variable) %>%
  # Calculating the stats
```

```

summarise(Mean = round(mean(Value, na.rm = TRUE), 2),
          "Standard Deviation" = round(sd(Value, na.rm = TRUE), 2),
          Minimum = round(min(Value, na.rm = TRUE), 2),
          Maximum = round(max(Value, na.rm = TRUE), 2)) %>%
kable()

# Visualizing response variables

hist_1 <- data %>%
  ggplot(aes(x = price)) +
  geom_histogram(fill = "burlywood2", color = "burlywood2", bins = 20) +
  labs(x = "Monthly Rent",
       y = "Frequency") +
  geom_vline(aes(xintercept = mean(price)),
             color = "black", linewidth = 0.5) +
  annotate("text",
          x = 2159 + 200,
          y = 200,
          label = "Mean: $2,159",
          color = "black",
          hjust = 0) +
  scale_x_continuous(labels = label_dollar()) +
  theme_classic()

hist_2 <- data %>%
  mutate(high_price = ifelse(high_price == 0,
                             "Rent Below\nAverage",
                             "Rent Above\nAverage")) %>%
  ggplot(aes(x = as.factor(high_price))) +
  geom_bar(fill = "grey") +
  labs(x = "", y = "Frequency") +
  geom_text(stat = "count", aes(label = after_stat(count)),
          vjust = 2, color = "white") +
  theme_classic()

hist_1 + hist_2 + plot_layout(widths = c(2, 1)) + plot_annotation(
  title = "Distributions of Toronto Monthly Apartment Rent and High Price Indicator (2018)"
)

# Making boxplots of price by bed/bath/den count

boxplot_1 <- data %>%
  mutate-bedroom = as.factor-bedroom)) %>%
  ggplot(aes(x = bedroom, y = price)) +
  geom_boxplot(fill = "burlywood2") +
  labs(y = "Monthly Rent", x = "# of Bedrooms") +
  scale_y_continuous(labels = label_dollar()) +
  theme_classic()

boxplot_2 <- data %>%
  mutate-bathroom = as.factor-bathroom)) %>%
  ggplot(aes(x = bathroom, y = price)) +
  geom_boxplot(fill = "burlywood2") +

```

```

labs(y = "Monthly Rent", x = "# of Bathrooms") +
scale_y_continuous(labels = label_dollar()) +
theme_classic()

boxplot_3 <- data %>%
  mutate(den = as.factor(den)) %>%
  ggplot(aes(x = den, y = price)) +
  geom_boxplot(fill = "burlywood2") +
  labs(y = "Monthly Rent", x = "# of Dens") +
  scale_y_continuous(labels = label_dollar()) +
  theme_classic()

boxplot_1 + boxplot_2 + boxplot_3 + plot_annotation(
  title = "Distributions of Toronto Rent Prices by Room Counts (2018)"
)

# Making a correlation matrix

par(font = 1)

corrplot(round(cor(select(data, -c(high_price, lat, long))), 2),
  method = "color",
  type = "upper",
  addCoef.col = "black",
  tl.col = "black",
  tl.srt = 30,
  tl.cex = 0.7,
  cl.ratio = 0.35,
  cl.length = 5,
  cl.cex = 0.7,
  number.cex = 0.7,
  col = colorRampPalette(c("#4477AA", "#77AADD", "#FFFFFF", "#EE9988", "#BB4444"))(200),
  mar = c(0, 0, 1, 0),
  diag = TRUE)

# Making frequency tables

data %>%
  tabyl(bedroom, high_price) %>%
  adorn_totals(c("row", "col")) %>%
  adorn_percentages("row") %>%
  adorn_pct_formatting(digits = 0) %>%
  adorn_ns() %>%
  kable(col.names = c("# of Bedrooms", "Rent Below Average", "Rent Above Average", "Total"))

data %>%
  tabyl(bathroom, high_price) %>%
  adorn_totals(c("row", "col")) %>%
  adorn_percentages("row") %>%
  adorn_pct_formatting(digits = 0) %>%
  adorn_ns() %>%
  kable(col.names = c("# of Bathrooms", "Rent Below Average", "Rent Above Average", "Total"))

```

```

# Linear multiplicative model

linear_mult <- lm(price ~ bedroom * bathroom, data = data)

tidy(linear_mult, conf.int = TRUE, conf.level = 0.95) %>%
  mutate_if(is.numeric, round, 2) %>%
  kable()

# Linear additive model

linear_add <- lm(price ~ bedroom + bathroom, data = data)

tidy(linear_add, conf.int = TRUE, conf.level = 0.95) %>%
  mutate(across(where(is.numeric) & !p.value, ~ round(.x, digits = 2))) %>%
  kable()

# Residual plot

plot(linear_add, which = 1)

# Q-Q plot

plot(linear_add, which = 2)

# Multiplicative logistic model

logistic_mult <- glm(high_price ~ bedroom * bathroom, data = data)

tidy(logistic_mult, conf.int = TRUE, conf.level = 0.95) %>%
  mutate_if(is.numeric, round, 2) %>%
  kable()

# Additive logistic model

logistic_add <- glm(high_price ~ bedroom + bathroom, data = data)

tidy(logistic_add, conf.int = TRUE, conf.level = 0.95) %>%
  mutate_if(is.numeric, round, 2) %>%
  kable()

tidy(logistic_add, conf.int = TRUE, conf.level = 0.95, exponentiate = TRUE) %>%
  mutate_if(is.numeric, round, 2) %>%
  kable()

# Binned residuals plot

plot(binned_residuals(logistic_mult)) +
  theme_classic()

```